

1.1 Assignment

a. Run Keras MNIST MLP Example

```
In [1]: '''Trains a simple deep NN on the MNIST dataset.
Gets to 98.40% test accuracy after 20 epochs
(there is *a lot* of margin for parameter tuning).
2 seconds per epoch on a K520 GPU.
'''

from tensorflow import keras
from tensorflow.keras.datasets import mnist
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout
from tensorflow.keras.optimizers import RMSprop

batch_size = 128
num_classes = 10
epochs = 20

# the data, split between train and test sets
(x_train, y_train), (x_test, y_test) = mnist.load_data()

x_train = x_train.reshape(60000, 784)
x_test = x_test.reshape(10000, 784)
x_train = x_train.astype('float32')
x_test = x_test.astype('float32')
x_train /= 255
x_test /= 255
print(x_train.shape[0], 'train samples')
print(x_test.shape[0], 'test samples')

# convert class vectors to binary class matrices
y_train = keras.utils.to_categorical(y_train, num_classes)
y_test = keras.utils.to_categorical(y_test, num_classes)

model = Sequential()
model.add(Dense(512, activation='relu', input_shape=(784,)))
model.add(Dropout(0.2))
model.add(Dense(512, activation='relu'))
model.add(Dropout(0.2))
model.add(Dense(num_classes, activation='softmax'))

model.summary()

model.compile(loss='categorical_crossentropy',
              optimizer=RMSprop(),
              metrics=['accuracy'])

history = model.fit(x_train, y_train,
                    batch_size=batch_size,
                    epochs=epochs,
                    verbose=1,
                    validation_data=(x_test, y_test))
score = model.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])
```

60000 train samples
 10000 test samples
 Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 512)	401920
dropout (Dropout)	(None, 512)	0
dense_1 (Dense)	(None, 512)	262656
dropout_1 (Dropout)	(None, 512)	0
dense_2 (Dense)	(None, 10)	5130

Total params: 669,706
 Trainable params: 669,706
 Non-trainable params: 0

Epoch 1/20

469/469 [=====] - 5s 10ms/step - loss: 0.2466 - accuracy: 0.9232 - val_loss: 0.1056 - val_accuracy: 0.9680

Epoch 2/20

469/469 [=====] - 4s 9ms/step - loss: 0.1058 - accuracy: 0.9687 - val_loss: 0.0757 - val_accuracy: 0.9784

Epoch 3/20

469/469 [=====] - 5s 10ms/step - loss: 0.0759 - accuracy: 0.9772 - val_loss: 0.0704 - val_accuracy: 0.9790

Epoch 4/20

469/469 [=====] - 5s 10ms/step - loss: 0.0611 - accuracy: 0.9815 - val_loss: 0.0814 - val_accuracy: 0.9777

Epoch 5/20

469/469 [=====] - 4s 9ms/step - loss: 0.0510 - accuracy: 0.9858 - val_loss: 0.0734 - val_accuracy: 0.9811

Epoch 6/20

469/469 [=====] - 5s 10ms/step - loss: 0.0448 - accuracy: 0.9863 - val_loss: 0.0686 - val_accuracy: 0.9829

Epoch 7/20

469/469 [=====] - 4s 9ms/step - loss: 0.0388 - accuracy: 0.9878 - val_loss: 0.0846 - val_accuracy: 0.9821

Epoch 8/20

469/469 [=====] - 4s 9ms/step - loss: 0.0345 - accuracy: 0.9898 - val_loss: 0.0784 - val_accuracy: 0.9817

Epoch 9/20

469/469 [=====] - 4s 9ms/step - loss: 0.0321 - accuracy: 0.9909 - val_loss: 0.0866 - val_accuracy: 0.9832

Epoch 10/20

469/469 [=====] - 4s 9ms/step - loss: 0.0284 - accuracy: 0.9918 - val_loss: 0.0882 - val_accuracy: 0.9829

Epoch 11/20

469/469 [=====] - 4s 9ms/step - loss: 0.0252 - accuracy: 0.9925 - val_loss: 0.0998 - val_accuracy: 0.9843

Epoch 12/20

469/469 [=====] - 4s 9ms/step - loss: 0.0246 - accuracy: 0.9927 - val_loss: 0.1062 - val_accuracy: 0.9834

Epoch 13/20

```
469/469 [=====] - 4s 9ms/step - loss: 0.0244 - accur
acy: 0.9933 - val_loss: 0.0997 - val_accuracy: 0.9833
Epoch 14/20
469/469 [=====] - 4s 9ms/step - loss: 0.0217 - accur
acy: 0.9935 - val_loss: 0.1026 - val_accuracy: 0.9841
Epoch 15/20
469/469 [=====] - 4s 9ms/step - loss: 0.0187 - accur
acy: 0.9949 - val_loss: 0.1001 - val_accuracy: 0.9841
Epoch 16/20
469/469 [=====] - 4s 9ms/step - loss: 0.0199 - accur
acy: 0.9946 - val_loss: 0.1196 - val_accuracy: 0.9835
Epoch 17/20
469/469 [=====] - 4s 9ms/step - loss: 0.0191 - accur
acy: 0.9946 - val_loss: 0.1169 - val_accuracy: 0.9824
Epoch 18/20
469/469 [=====] - 4s 9ms/step - loss: 0.0195 - accur
acy: 0.9948 - val_loss: 0.1182 - val_accuracy: 0.9848
Epoch 19/20
469/469 [=====] - 4s 9ms/step - loss: 0.0186 - accur
acy: 0.9948 - val_loss: 0.1281 - val_accuracy: 0.9832
Epoch 20/20
469/469 [=====] - 4s 9ms/step - loss: 0.0164 - accur
acy: 0.9955 - val_loss: 0.1207 - val_accuracy: 0.9840
Test loss: 0.12066762894392014
Test accuracy: 0.984000027179718
```

b. Run PySpark Example

```

In [5]: #
# Licensed to the Apache Software Foundation (ASF) under one or more
# contributor license agreements. See the NOTICE file distributed with
# this work for additional information regarding copyright ownership.
# The ASF licenses this file to You under the Apache License, Version 2.0
# (the "License"); you may not use this file except in compliance with
# the License. You may obtain a copy of the License at
#
# http://www.apache.org/licenses/LICENSE-2.0
#
# Unless required by applicable law or agreed to in writing, software
# distributed under the License is distributed on an "AS IS" BASIS,
# WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied.
# See the License for the specific language governing permissions and
# limitations under the License.
#

import sys
from random import random
from operator import add

from pyspark.sql import SparkSession

if __name__ == "__main__":
    """
    Usage: pi [partitions]
    """
    spark = SparkSession\
        .builder\
        .appName("PythonPi")\
        .getOrCreate()

    partitions = int(sys.argv[1]) if len(sys.argv) > 1 & sys.argv[1].isdigit() == 1
    n = 100000 * partitions

    def f(_):
        x = random() * 2 - 1
        y = random() * 2 - 1
        return 1 if x ** 2 + y ** 2 <= 1 else 0

    count = spark.sparkContext.parallelize(range(1, n + 1), partitions).map(f).reduce(
        print("Pi is roughly %f" % (4.0 * count / n))

    spark.stop()

```

Pi is roughly 3.134960

I added one logic to handle the partitions logic since argument is coming as String.

In []:

